

ABC Company

Services Catalog — Comprehensive Guide to Our Technology Solutions

Document ID:	ABC-DOC-002
Version:	v1.0
Date:	June 2026
Prepared By:	ABC Company — Business Development
Classification:	Business Confidential

ABC Company — Colombo | London | Melbourne

Services Overview

ABC Company provides end-to-end technology services across nine practice areas. The table below summarises each service, its core focus, the primary technologies involved, and a typical project timeline. Detailed descriptions of each service, including team composition, delivery approach, and decision frameworks, follow in subsequent sections of this catalog.

Service	Core Focus	Key Technologies	Typical Timeline
Custom Web Application Development	End-to-end bespoke web applications	React, Next.js, Node.js/NestJS, Python/FastAPI	2–6 months
Mobile App Development	Native and cross-platform iOS & Android apps	React Native, Flutter, Swift, Kotlin	3–8 months
Cloud Infrastructure & DevOps	Cloud architecture, migration, and CI/CD automation	AWS, Azure, GCP, Kubernetes, Terraform	2 weeks – 6 months
AI/ML & Data Engineering	LLM applications, RAG systems, data pipelines	OpenAI, Anthropic, LangChain, PyTorch, Airflow	4 weeks – 6 months
UI/UX Design	User research, interface design, design systems	Figma, Adobe Creative Suite, Maze, Hotjar	2–6 weeks
Quality Assurance & Testing	Manual, automated, performance, and security testing	Playwright, Cypress, Jest, pytest, k6, OWASP ZAP	Ongoing (sprint-embedded)
Staff Augmentation & Dedicated Teams	Skilled engineers embedded in client teams	Flexible model: individual, pod, or full team	3-month minimum
Maintenance, Support & Managed Services	Post-launch support, monitoring, enhancements	Full-stack toolchain matching client environment	Ongoing SLA-based
Product Strategy & Technical Consulting	CTO-as-a-Service, due diligence, architecture review	Methodology-agnostic; tool-agnostic advisory	Weekly sessions or project-based

All engagements begin with a scoping call to understand your requirements, constraints, and success criteria. We provide fixed-price quotes for well-defined scopes and time-and-materials arrangements for exploratory engagements. Contact our Business Development team at **bd@abccompany.io** to schedule a consultation.

Service 1: Custom Web Application Development

Core Technologies	React, Next.js, Angular, Vue.js, Node.js/NestJS, Python/FastAPI/Django
Typical Team	1 Tech Lead, 2–4 Frontend Devs, 1–3 Backend Devs, 1 QA Engineer, 1 DevOps (part-time)
Typical Timeline	2–6 months depending on complexity
Deliverables	Responsive web apps, PWAs, SPAs, SSR applications, admin dashboards, customer portals

Custom web application development sits at the heart of ABC Company's engineering practice. We have built and deployed over one hundred production web applications across industries including financial services, healthcare, logistics, retail, and professional services. Our approach is deliberately end-to-end: we engage from the initial concept and requirements workshop through architecture design, iterative development, QA, and into production deployment. Clients who come to us typically have a specific problem that cannot be adequately solved by an off-the-shelf SaaS product — they need something tailored precisely to their business logic, data model, and user workflows.

Our development methodology is component-driven and API-first. On the frontend we default to React for its ecosystem maturity, component composability, and the breadth of talent available. For projects where SEO is a primary concern, or where server-side rendering yields meaningful performance improvements for end users, we use Next.js with its hybrid static and server-side rendering capabilities. Angular remains our choice for large enterprise applications where the opinionated framework structure, built-in dependency injection, and TypeScript-first design reduce long-term maintenance costs on teams of ten or more developers. Vue.js enters the conversation for projects requiring a gentler learning curve or for teams that already have Vue expertise in-house.

On the backend, our two primary choices are Node.js with NestJS and Python with FastAPI or Django. NestJS is selected for projects requiring real-time features such as WebSocket-heavy collaborative tools, high-throughput API gateways, or teams that want end-to-end TypeScript from browser to server. FastAPI is our default for data-intensive applications, ML-adjacent services, or when the broader Python ecosystem (pandas, SQLAlchemy, Celery) maps cleanly onto the problem. Django with Django REST Framework remains the pragmatic choice for content-rich applications with complex permission systems or when rapid admin interface generation is operationally valuable.

Performance is designed in from day one rather than retrofitted. We establish Core Web Vitals targets — Largest Contentful Paint under 2.5 seconds, Cumulative Layout Shift below 0.1, and Interaction to Next Paint under 200 milliseconds — as acceptance criteria alongside functional requirements. Code splitting, lazy loading, image optimisation through modern formats such as WebP and AVIF, aggressive CDN caching strategies, and database query profiling are standard practice across all engagements. We run Lighthouse audits on every pull request as part of our CI pipeline to catch performance regressions before they reach production.

Accessibility is treated as a non-negotiable engineering requirement, not an afterthought. All applications we deliver target WCAG 2.1 Level AA compliance as a minimum, with AA+ for public sector and healthcare clients where legal obligations apply. Our engineers use axe-core integrated into the automated test suite for continuous accessibility scanning, supplemented by manual testing with NVDA, VoiceOver, and JAWS screen readers. Keyboard navigability, appropriate ARIA labelling, sufficient colour contrast ratios, and focus management in single-page application route transitions are reviewed explicitly during our accessibility-specific sprint ceremonies.

Responsive design is delivered through a mobile-first methodology. We design and build for the smallest target viewport first, then progressively enhance the layout for larger breakpoints. This discipline ensures that the core user journey is always fast and functional on constrained devices, and that desktop enhancements are genuinely additive rather than making mobile an afterthought. Progressive Web App capabilities — service workers, Web App Manifest,

offline caching strategies — are added where they provide tangible user value, particularly for applications used in connectivity-limited environments such as field service tools or warehouse management systems.

A typical engagement is staffed with one Tech Lead who owns the architecture decisions and maintains technical quality across the team, two to four Frontend Developers depending on UI complexity, one to three Backend Developers, one QA Engineer embedded throughout development rather than engaged only at the end, and a part-time DevOps Engineer who provisions environments, builds the CI/CD pipeline, and manages infrastructure as code from day one. Timeline expectations range from two months for focused single-module applications to six months for multi-role platforms with complex integrations.

When to Choose This Service

You need a custom web application that does not fit off-the-shelf solutions. You require specific integrations, high performance, custom business logic, or a user experience that standard SaaS products cannot deliver.

Service 2: Mobile App Development

Core Technologies	React Native with Expo (primary), Flutter, Swift (iOS native), Kotlin (Android native)
Typical Team	1 Mobile Lead, 2–3 Mobile Developers, 1 Backend Developer, 1 QA (device testing)
Typical Timeline	3–8 months
Deliverables	Consumer apps, enterprise mobile solutions, offline-first apps, IoT companion apps

Mobile applications require a fundamentally different engineering mindset to web applications. Users interact with phones in short, high-intent bursts; connectivity is unreliable; device hardware — camera, GPS, biometrics, accelerometer — is a feature, not an edge case; and app store distribution adds a layer of governance and review that demands disciplined release engineering. ABC Company has delivered mobile applications across consumer, enterprise, and field operations contexts, accumulating deep experience across the iOS and Android ecosystems and the leading cross-platform frameworks.

The first architectural decision in any mobile project is whether to invest in a cross-platform codebase or build separate native applications. Our guidance is grounded in honest assessment rather than default preference. Cross-platform with React Native is the right choice for the majority of business applications: it delivers near-native performance, allows a single engineering team to serve both platforms, and shares business logic and most UI components. React Native with Expo is our primary framework because Expo's managed workflow dramatically simplifies build infrastructure, over-the-air updates, push notification handling, and app store submission workflows. Flutter is a compelling alternative when pixel-perfect design fidelity is paramount — its Skia-based rendering engine bypasses native UI components entirely, giving designers precise control at the cost of a larger binary size and a smaller but rapidly growing ecosystem.

Native Swift for iOS and Kotlin for Android are recommended when the application is deeply integrated with platform-specific frameworks — ARKit and RealityKit for augmented reality, HealthKit for clinical data, or intensive CoreML on-device inference on iOS; Jetpack Compose animations, WorkManager background processing, or Android Enterprise management APIs on Android. Native development also becomes relevant when the existing engineering team has strong platform-specific expertise and the long-term maintenance team will remain platform-specialized.

Mobile UX patterns diverge significantly from web patterns. Thumb zones, bottom navigation bars, swipe gestures, bottom sheets for secondary actions, and haptic feedback are conventions that users expect; violating them creates friction that reduces retention. Our mobile designers work within Apple Human Interface Guidelines and Google Material Design 3 principles as baselines, adapting where brand identity requires divergence but never abandoning the underlying interaction principles. Onboarding flows receive particular attention: research consistently shows that the first three minutes of a new user's experience determines 30-day retention.

Offline-first architecture is a differentiating capability for enterprise mobile applications. We implement local SQLite or Realm databases as the single source of truth, with a synchronisation layer that queues mutations made offline and resolves conflicts on reconnection. Field service applications, inventory management tools, and inspection checklists are typical use cases where this pattern converts a connectivity-dependent application into one that works reliably regardless of signal. Push notifications, biometric authentication via Face ID and fingerprint, deep linking, and camera-based barcode scanning are capabilities we integrate routinely.

App store submission is a process that surprises many clients encountering it for the first time. Apple's App Store review process has a median time of 24–48 hours but can exceed a week for applications using certain entitlements. We front-load App Store Connect configuration, privacy manifest requirements, export compliance declarations, and TestFlight beta distribution in the first sprint so that the submission pipeline is tested well before the planned release

date. Google Play's review process is generally faster but requires similar rigour around Data Safety section declarations and target API level compliance.

A typical mobile engagement is staffed with one Mobile Lead who makes framework and architecture decisions, two to three Mobile Developers, one Backend Developer to build or extend the APIs the mobile clients consume, and one QA Engineer with access to a device farm covering the top ten Android devices by market share in the target geography plus the current and immediately preceding generations of iPhone. Timelines run from three months for a focused MVP to eight months for feature-rich consumer applications.

When to Choose This Service

Your users primarily interact via smartphones. You need native-quality mobile experiences, push notifications, biometric authentication, offline capability, or device hardware access such as camera, GPS, or NFC.

Service 3: Cloud Infrastructure & DevOps

Core Technologies	AWS (primary), Azure, GCP, Kubernetes, Docker, Terraform, GitHub Actions, GitLab CI, Jenkins
Typical Team	1 Cloud Architect, 1–2 DevOps Engineers
Typical Timeline	2–8 weeks for pipeline setup; 2–6 months for cloud migration
Deliverables	Cloud architecture design, CI/CD pipelines, infrastructure as code, monitoring, cost optimisation

Modern software products live and die by their infrastructure. An application that crashes under load, deploys unpredictably, or accumulates cloud costs without visibility will undermine even the best-engineered features. ABC Company's Cloud Infrastructure and DevOps practice provides the engineering discipline to build cloud environments that are reliable, secure, cost-efficient, and automated from day one. We hold twelve AWS certifications across our team including Solutions Architect Professional, DevOps Engineer Professional, and Security Specialty, and maintain active practice areas in Microsoft Azure and Google Cloud Platform.

Infrastructure as Code is not optional for us — it is the starting point. Every cloud resource we provision is defined in Terraform, with state managed remotely in S3 with DynamoDB locking on AWS, or equivalent backends on Azure and GCP. Environments — development, staging, and production — are identical by construction, eliminating the class of production incidents caused by configuration drift. Terraform modules are version-pinned and reviewed before merging to main, and a Terraform plan is generated and reviewed as a mandatory step in every infrastructure pull request.

CI/CD maturity directly correlates with engineering velocity and production stability. Our standard pipeline architecture progresses through automated unit and integration test execution, static analysis and security scanning with tools such as Checkov for infrastructure and Semgrep for application code, Docker image build and push to a private registry, environment promotion through staging and then production with automated smoke tests at each gate, and structured rollback procedures triggered by health check failures. GitHub Actions is our default pipeline orchestrator for its native integration with repository events; GitLab CI and Jenkins are supported for clients with existing toolchain investments.


Kubernetes is recommended for applications requiring fine-grained resource scheduling, sophisticated traffic management, or multi-tenant isolation. We deploy to Amazon EKS, Azure AKS, or Google GKE, managing cluster configuration through Helm charts and GitOps workflows using Argo CD. For applications that do not need the full complexity of Kubernetes, we recommend AWS ECS with Fargate or equivalent serverless container platforms, which reduce operational overhead significantly while preserving the benefits of container-based deployment.

Our multi-cloud strategy acknowledges that vendor lock-in carries real risk. We architect for portability by abstracting provider-specific services behind interface boundaries, using open standards such as OpenTelemetry for observability, and ensuring that database choices are not tied to proprietary hosted variants where alternatives exist. In practice, most clients operate primarily on a single cloud provider and use a secondary provider for disaster recovery or data sovereignty compliance rather than active-active multi-cloud — our advisory accounts for this operational reality.

Cost optimisation is delivered through a combination of architectural right-sizing, Reserved Instance and Savings Plan coverage for predictable baseline workloads, spot or preemptible instances for fault-tolerant batch and background processing, and tagging strategies that enable cost allocation by team, product, and environment. We establish cost

monitoring dashboards with budget alerts from the first week of every engagement. Clients typically see 20–40% cloud spend reduction following a dedicated optimisation sprint.

Security hardening of cloud infrastructure follows CIS Benchmarks as a baseline. We implement least-privilege IAM policies, VPC network segmentation with private subnets for data tier resources, encryption at rest and in transit as standard, secrets management through AWS Secrets Manager or HashiCorp Vault, and CloudTrail or equivalent audit logging enabled across all accounts. GuardDuty or equivalent threat detection is enabled from day one, not added as a compliance afterthought.



Service 4: AI/ML & Data Engineering

Core Technologies	OpenAI GPT-4o, Anthropic Claude, LangChain, TensorFlow, PyTorch, scikit-learn, Hugging Face, Apache Airflow, dbt, Snowflake
Typical Team	1 ML Engineer/Data Scientist, 1–2 Backend Developers, 1 Data Engineer
Typical Timeline	4–12 weeks for AI features; 2–6 months for full data platform
Deliverables	LLM applications, RAG systems, chatbots, recommendation engines, predictive analytics, data pipelines, ETL/ELT, data warehousing

The rapid maturation of large language model APIs and the surrounding tooling ecosystem has created genuine commercial opportunity for organisations that can move quickly and responsibly. ABC Company's AI/ML and Data Engineering practice spans the full spectrum from embedding LLM capabilities into existing products through to building purpose-built data platforms that transform raw operational data into decision-grade intelligence. We work with the frontier model APIs from OpenAI and Anthropic, open-source models deployed on private infrastructure for data-sensitive use cases, and the growing ecosystem of tooling built around them.

Retrieval-Augmented Generation is our most frequently delivered AI architecture pattern. RAG addresses the fundamental limitation of LLMs — their knowledge is frozen at training time and does not include your organisation's specific documents, policies, product data, or customer records. A well-architected RAG system ingests your knowledge base into a vector store, encodes queries as embeddings, retrieves semantically relevant chunks at inference time, and injects them into the LLM prompt context. The result is an AI system that speaks accurately about your domain without the cost, risk, or latency of fine-tuning. We have delivered RAG systems using Qdrant, Pinecone, pgvector, and Weaviate, with ingestion pipelines handling PDFs, Word documents, web scrapes, and structured database exports.

Fine-tuning is appropriate in a narrower set of circumstances than marketing materials suggest. It makes sense when the task requires a very specific output format, tone, or domain vocabulary that cannot be achieved through prompt engineering; when latency or cost requirements demand a smaller model than the frontier options; or when data sovereignty requirements prohibit sending inputs to external API endpoints. We use supervised fine-tuning via OpenAI's fine-tuning API for GPT-4o-mini use cases, and Parameter-Efficient Fine-Tuning techniques such as LoRA and QLoRA for open-source base models deployed on our clients' own GPU infrastructure.

Data engineering underpins every AI and analytics capability. Raw data in operational systems — transactional databases, SaaS APIs, event streams — must be extracted, transformed, quality-checked, and loaded into analytical stores before models or dashboards can consume it reliably. We build these pipelines using Apache Airflow for orchestration, dbt for SQL-based transformation with full lineage tracking and testing, and Snowflake or BigQuery as the analytical warehouse depending on existing client toolchain commitments. For high-velocity event streams we introduce Kafka with Flink or Spark Structured Streaming for real-time processing requirements.

Responsible AI is not a compliance checkbox for us — it is an engineering discipline embedded in how we design and test AI features. We apply evaluation frameworks that measure factual accuracy, hallucination rate, bias in outputs across demographic dimensions, and prompt injection resilience before any AI feature reaches production. Human-in-the-loop review gates are designed for high-stakes decisions. Data provenance is maintained so that the sources underlying any AI-generated output can be audited. We engage with our clients' legal and compliance teams early in AI projects to ensure that regulatory requirements — including the EU AI Act where applicable — are designed into the system rather than bolted on.

A typical AI feature engagement is staffed with one ML Engineer or Data Scientist who owns model selection, evaluation, and RAG architecture, one to two Backend Developers who build the APIs and integrations around the AI components, and one Data Engineer for projects requiring substantial data pipeline work. Full data platform engagements also include a Data Architect in the initial design phase. Timelines run from four weeks for a focused LLM feature integration to six months for a comprehensive data platform with multiple downstream applications.

When to Choose This Service

You want to embed AI capabilities into an existing product, build a purpose-built AI application, or transform raw operational data into decision-grade analytics. We provide honest assessment of where AI creates genuine value versus where simpler solutions are more appropriate.

Service 5: UI/UX Design

Core Tools	Figma (primary), Adobe Creative Suite, Maze (user testing), Hotjar (analytics)
Typical Team	1–2 UI/UX Designers, 1 UX Researcher (for larger projects)
Typical Timeline	2–6 weeks for design phase; ongoing for design system maintenance
Deliverables	User research & personas, wireframes & prototypes, high-fidelity UI designs, design systems, usability testing reports

Great software is both functional and genuinely pleasurable to use. The gap between software that works and software that users actually adopt is almost always a design gap. ABC Company's UI/UX Design practice operates as a full-spectrum design service: we conduct user research to validate assumptions about how real users think and behave, translate that research into information architectures and interaction designs, produce high-fidelity visual designs that balance brand expression with usability principles, and build the design systems that allow engineering teams to implement consistently at speed.

Our design process is grounded in Design Thinking: Empathise, Define, Ideate, Prototype, and Test. We run user interviews and contextual inquiry sessions to develop genuine empathy for the target users' tasks, frustrations, and mental models. Synthesis sessions produce user personas and journey maps that the whole product team references throughout development. Problem definition is made explicit: we write 'How Might We' statements that frame the design challenge before generating solutions, preventing the common failure mode of designing solutions to the wrong problem.

We apply Atomic Design principles to build interfaces from the smallest reusable elements up. Atoms — buttons, input fields, labels, icons — are designed with all states (default, hover, focus, disabled, error) before being composed into molecules and organisms. This discipline means that when a designer hands off to an engineer, there is no ambiguity about interactive states, spacing tokens, or typography scales. The Figma component library mirrors the React component library: every Figma component has a direct counterpart in the frontend codebase, and the design system documentation in Figma is the source of truth for both designers and engineers.

Accessibility is integrated into the design process from the first wireframe. Colour combinations are checked for WCAG AA contrast ratios in Figma using the Contrast plugin before any mockup is approved. Interactive component designs include keyboard navigation flows and screen reader annotation layers. We conduct usability testing sessions that explicitly recruit participants with disabilities, using assistive technology in realistic conditions, and incorporate findings into design iterations before engineering begins.

Usability testing is conducted in two modes: moderated sessions for depth and unmoderated sessions at scale. Moderated sessions use Maze or UserTesting platforms to observe participants completing task scenarios with prototypes, capturing think-aloud commentary, navigation paths, and time-on-task metrics. Unmoderated studies with larger sample sizes validate quantitative metrics such as task completion rate and System Usability Scale scores. For live products we integrate Hotjar for session recording and heatmap analysis to ground iterative design decisions in real behaviour data.

Design handoff is a process we take seriously because poor handoff is where design fidelity is lost. We use Figma's developer mode to expose spacing, typography, and colour values directly to engineers, supplemented by written component specifications for complex interactions. We run handoff review sessions with the engineering team for each major feature, walking through component behaviour, edge cases, empty states, and error states. Designers remain accessible throughout development to answer questions and review implementations, preventing the drift that occurs when design and engineering become isolated phases.

Service 6: Quality Assurance & Testing

Core Technologies	Playwright (E2E), Cypress, Jest, pytest, k6 (performance), OWASP ZAP (security), Appium (mobile)
Typical Team	1–2 QA Engineers, 1 Performance Engineer (if needed)
Typical Timeline	Ongoing — embedded in development sprints
Deliverables	Test strategy, automated test suites, performance benchmarks, security audit reports, CI-integrated pipelines

Quality is not a phase that happens at the end of development — it is a continuous discipline woven through every sprint. ABC Company's Quality Assurance and Testing practice is built on the principle of shifting quality left: the earlier a defect is detected, the cheaper it is to fix. A bug found by a developer writing a unit test costs minutes to resolve. The same bug found by a customer in production costs hours of engineering time, potential data correction, and real damage to user trust. Our QA engineers are embedded within development teams from sprint one, not brought in at the end to validate a finished product.

Our testing strategy follows the test pyramid model. The foundation is a comprehensive unit test suite, targeting 80% code coverage on business logic and utility functions with Jest for JavaScript/TypeScript and pytest for Python. Unit tests run in under thirty seconds on every commit, providing immediate feedback to developers. The middle layer is integration testing: API contract tests using pytest with an in-memory database, component integration tests in React Testing Library that test components with their real dependencies rather than excessive mocking, and service-to-service contract tests using Pact for microservice architectures. The apex of the pyramid is end-to-end testing with Playwright, which is our primary E2E framework for its superior reliability, cross-browser support including Chromium, Firefox, and WebKit, and first-class support for modern web application patterns including shadow DOM and web components.

Performance testing is conducted using k6, which allows us to express load test scenarios as JavaScript code, version them in the same repository as the application, and execute them as part of the CI pipeline on a dedicated performance testing environment. We establish baseline performance benchmarks early in the project — typically at the first integration of a major API endpoint — and test against those baselines in every deployment pipeline run. Stress testing, spike testing, and soak testing are conducted before significant traffic events such as product launches or marketing campaigns. Results are presented in k6 Cloud or Grafana dashboards with clear pass/fail thresholds for p95 response time, error rate, and throughput.

Security testing is integrated at multiple levels. OWASP ZAP is run in automated scan mode against every staging deployment, checking for the OWASP Top Ten vulnerability classes including SQL injection, XSS, CSRF, and insecure direct object references. Dependency vulnerability scanning via Dependabot and npm audit runs on every pull request. For applications handling sensitive data or subject to compliance frameworks, we conduct structured manual penetration testing at major release milestones. Infrastructure security scanning using Checkov and tfsec is applied to all Terraform code before plan execution.

Mobile application testing requires device-specific considerations beyond what emulators capture reliably. We maintain access to physical device testing through BrowserStack App Automate for broad coverage, supplementing with devices in our own lab for the highest-priority form factors. Appium drives automated UI tests on native iOS and Android applications; for React Native and Flutter applications we integrate framework-native testing utilities. OS version compatibility, screen size adaptation, permission grant and denial flows, and interruption handling — incoming calls, notifications — are explicitly tested.

The cultural dimension of QA is as important as the tooling. We advocate for a Definition of Done that includes passing automated tests, no new critical or high bugs, and peer code review as non-negotiable gates before any story is closed. Bug bash sessions before major releases bring together the whole team — including product managers and designers — for structured exploratory testing. Post-release, we monitor production error rates via Sentry or equivalent real-user monitoring tools and treat any regression in error rate as a P1 engineering priority.

Service 7: Staff Augmentation & Dedicated Teams

Models	Individual Contributor, Pod (3–5 engineers + lead), Full Team (6+ including PM, QA, DevOps)
Flexibility	Scale up/down with 2-week notice; minimum 3-month engagement
Onboarding	Productive within 1 week
Vetting	4-stage assessment; <10% acceptance rate

Staff Augmentation and Dedicated Teams is the service we are most frequently asked about by clients who have an established engineering function but face a gap: a skills shortage in a specific technology, a temporary capacity surge driven by an accelerated roadmap, or a need for senior technical leadership during a critical transition. We place engineers who are ABC Company employees — fully vetted, actively supported, and professionally developed within our organisation — into the client's team. They attend standups, contribute to the same codebase, use the same tools, and are indistinguishable from direct hires in practice.

Our vetting process is rigorous because our reputation is staked on every engineer we place. Candidates complete a four-stage assessment: an initial technical screen covering fundamentals in their primary stack, a take-home project assessed on code quality, testing approach, and documentation, a live technical interview with two senior engineers covering architecture decisions and problem-solving approach, and a professional reference check. Fewer than ten percent of applicants complete this process successfully. The result is a talent pool where every engineer has been validated to a consistent bar — clients are not conducting their own technical screening of candidates we forward; they are selecting from a pre-validated pool.

Speed of contribution is a priority our clients consistently name. We target productivity within the first week of an engagement. This is achieved through structured pre-engagement preparation: the assigned engineer receives detailed onboarding materials about the client's domain, technical stack, code conventions, and team norms before day one. ABC Company maintains a shared knowledge base of common technology stacks and integration patterns that engineers use to reduce the learning curve on familiar technologies. We assign an ABC Company engagement lead to monitor onboarding progress and remove obstacles during the critical first month.

We offer three staffing models to match different client needs. The Individual Contributor model places one engineer with specialised skills — a senior React Native developer, a Kubernetes specialist, a data engineer with Airflow expertise — directly into the client's existing team. The Pod model provides three to five engineers plus a technical lead who manages the group, suited for clients who want a self-managing delivery unit that can take full ownership of a product area or module. The Full Team model provides six or more engineers including a dedicated Project Manager, QA Engineer, and DevOps Engineer, appropriate for clients who want complete product delivery accountability without building an in-house function.

Engagements require a minimum three-month commitment, which ensures that the ramp-up cost is amortised over a meaningful delivery period. Beyond that minimum, clients can scale the team up or down with two weeks' notice, allowing the engineering capacity to track roadmap priorities without the administrative burden of repeated hiring and redundancy processes. When a client's needs evolve — for example, shifting from a development-heavy phase to a maintenance phase — we adjust the skill composition of the team accordingly. ABC Company absorbs the internal transition cost.

The fundamental distinction between staff augmentation and traditional project outsourcing is accountability and integration. In a classic outsourced engagement, a vendor team operates in parallel to the client team with a defined scope, a fixed contract, and a handover at the end. Knowledge transfers are formal events; the client is dependent on the vendor's interpretation of requirements. In our augmentation model, engineers are embedded in the client's team,

operate under the client's processes, and build knowledge of the client's system that stays within the client's organisation. This fundamentally changes the risk profile and the alignment of incentives.

<10% CANDIDATE ACCEPTANCE RATE	1 week TIME TO CONTRIBUTION	2 weeks SCALE CHANGE NOTICE
---	---	---

Service 8: Maintenance, Support & Managed Services

Tiers	Basic, Standard, Premium
Response SLA	4 hours (Premium) — 48 hours (Basic)
Coverage	Business hours (Basic) — 24/7 (Premium)
Included	Bug fixes, security patches, dependency updates, performance monitoring, minor enhancements

Launching a product is the beginning of the engineering journey, not the end. Applications require ongoing care: security vulnerabilities in dependencies are disclosed continuously; frameworks release new versions; infrastructure costs shift as usage patterns evolve; users discover edge cases that automated testing did not anticipate; and businesses evolve requirements faster than any initial roadmap projected. ABC Company's Maintenance, Support, and Managed Services offering provides the ongoing engineering partnership that allows client products to remain healthy, secure, and improving without requiring the client to maintain a full in-house engineering team.

Our approach to support is proactive rather than purely reactive. Every managed application is monitored using a combination of application performance monitoring through New Relic or Datadog, error tracking through Sentry, infrastructure health through cloud-native monitoring services, and uptime monitoring through external probes. Alerting thresholds are established collaboratively with the client: we alert on error rate spikes, P95 latency degradation, disk and memory utilisation approaching thresholds, failed background job rates, and certificate expiry warnings. The goal is that our team detects and begins investigating an incident before the client or their customers report it.

Service Tier Comparison

Tier	Hours/Month	Coverage	Response SLA	Health Reports	Price Range
Basic	20 hours/month	Business hours	48 hours	No	\$3,000 – \$5,000 / month
Standard	40 hours/month	Extended hours (8am–10pm)	24 hours	Monthly	\$6,000 – \$10,000 / month
Premium	80+ hours/month	24/7	4 hours	Monthly + quarterly strategy	\$12,000 – \$20,000 / month

Monthly health reports are included in Standard and Premium tiers. Each report covers: application performance trends versus the prior month, error rate analysis with root cause notes for any spikes, dependency audit results listing packages with available security patches, infrastructure cost summary versus budget, planned maintenance activities for the coming month, and a running log of completed work items. These reports serve both operational and executive audiences: they provide the product team with actionable engineering visibility and give leadership a concise health indicator that does not require technical expertise.

Incident response follows a defined escalation path that varies by tier. For Premium tier clients with a four-hour SLA, on-call engineers are paged immediately on any P1 alert. A war room communication channel is established within fifteen minutes of a confirmed incident. Status updates are posted on a cadence defined at onboarding — typically every thirty minutes during active investigation. Post-incident reviews are conducted within forty-eight hours of

resolution, producing a written post-mortem that covers timeline, root cause, contributing factors, and remediation actions with owners and deadlines.

Managed services hours cover bug fixes and regression resolution, security patches and dependency updates, performance monitoring and optimisation, minor feature enhancements within the existing architecture, infrastructure maintenance, and documentation updates. Major new features requiring architecture changes, platform migrations, and full redesigns are scoped and priced separately as development engagements. This boundary is defined clearly in the service agreement and reviewed monthly to prevent ambiguity about what falls within the retainer.

Service 9: Product Strategy & Technical Consulting

Offerings	CTO-as-a-Service, Technical Due Diligence, Product Roadmap Workshops, Architecture Review
Typical Team	1 Senior Architect or CTO-level consultant
Engagement	Weekly advisory sessions, project-based assessments, or embedded part-time CTO
Rate	\$150–\$200 per hour

Not every engagement begins with a development sprint. Sometimes the most valuable thing ABC Company can provide is a clear-eyed external perspective on what to build, how to build it, or whether the system you have already built is fit for purpose. Our Product Strategy and Technical Consulting practice provides senior-level advisory services that give founders, product leaders, and investors the technical clarity they need to make high-stakes decisions with confidence.

CTO-as-a-Service is designed for early-stage companies that have a compelling product idea and a development team — whether in-house or outsourced — but lack a senior technical leader to set the architecture direction, make technology choices, manage the engineering team, and represent the technical function to investors and the board. A fractional CTO from ABC Company engages on a part-time basis — typically one to two days per week — providing the strategic technical leadership without the cost or commitment of a full-time executive hire. The engagement typically spans six to eighteen months, covering the period from initial product architecture through Series A or the point at which the company is ready to hire a permanent CTO. Deliverables include a documented technical strategy, architecture decision records, engineering team hiring criteria, and a technical roadmap that connects product ambitions to engineering milestones.

Technical Due Diligence is conducted for investors considering a stake in a technology company, and for acquirers evaluating a target. The assessment covers five dimensions: codebase quality (architecture, test coverage, technical debt, security vulnerabilities); infrastructure maturity (deployment automation, observability, disaster recovery); team capability (size relative to product complexity, key person risk, engineering process maturity); scalability headroom (current load versus designed capacity, identified bottlenecks); and IP ownership (open-source licence compliance, contractor IP assignment). The engagement typically takes two to three weeks and produces a written report with an executive summary, a detailed findings section graded by risk level, and a remediation roadmap with estimated effort. We maintain strict confidentiality throughout, working under NDA, and never share findings beyond the commissioning party.

Product Roadmap Workshops are structured facilitated sessions that bring together product, commercial, and engineering stakeholders to build alignment on priorities. A typical engagement runs as two full-day workshops bookending a two-week preparation phase. The preparation phase involves stakeholder interviews, review of existing product analytics and user research, competitive landscape analysis, and pre-reading materials distributed to all participants. The first workshop uses techniques from Impact Mapping and Story Mapping to connect business goals to user needs to potential features. The second workshop applies effort estimation, dependency mapping, and risk assessment to produce a prioritised, sequenced roadmap for the next two to four quarters. The output is a living Notion or Confluence document — not a static slide deck — that the team continues to update.

Architecture Review is the right engagement when an engineering team has moved fast and is beginning to feel the consequences: deployment is slow and fragile, the codebase is hard to navigate for new joiners, certain features require disproportionate effort to build, or the system is approaching the limits of its current scaling approach. Our review process begins with a structured interview programme across the engineering team, followed by a week of

code and infrastructure examination. We assess against the Well-Architected Framework for cloud infrastructure and against established software architecture principles — separation of concerns, bounded contexts, single responsibility, observable systems — for application code. The output is a findings report with prioritised recommendations categorised as immediate action items, medium-term refactoring investments, and long-term architectural evolution goals.

The right time to engage consulting rather than development is when the decision ahead of you is more valuable than the code. If you are deciding between three architecture approaches for a system that will run for the next five years, spending a week with a senior architect to evaluate the trade-offs is better economics than committing your development team to the wrong choice for three months. If you are raising a Series A and an investor has requested technical due diligence, a rigorous external assessment protects both parties and often accelerates the deal rather than slowing it. If your engineering team is consistently missing sprint commitments and you cannot identify why, an architecture review and team process assessment will surface root causes faster than incremental trial and error.

Consulting Engagement Formats

CTO-as-a-Service: Part-time technical leadership for companies without a permanent CTO. Typically 1–2 days per week.

Technical Due Diligence: 2–3 week assessment producing a written report with risk-graded findings and remediation roadmap.

Roadmap Workshop: Two facilitated full-day sessions plus two weeks of preparation. Output: prioritised two-to-four quarter roadmap.

Architecture Review: One-week code and infrastructure assessment producing prioritised recommendations across three time horizons.

Engaging ABC Company

Every engagement begins with a discovery conversation. We take time to understand your business context, technical landscape, team structure, and the outcome you are trying to achieve before proposing an approach or a price. We do not believe in one-size-fits-all proposals: the right team size, technology choices, and engagement model depend entirely on your specific situation.

Our Business Development team is available for an initial consultation at no charge. Bring your challenge, your constraints, and your timeline — we will be direct about what is achievable and honest about where we can and cannot add value.

Email	bd@abccompany.io
Website	www.abccompany.io
Colombo Office	Level 14, World Trade Centre, Echelon Square, Colombo 01, Sri Lanka
London Office	1 Canada Square, Canary Wharf, London E14 5AB, United Kingdom
Melbourne Office	Collins Street Tower, 447 Collins St, Melbourne VIC 3000, Australia

This document is prepared by ABC Company Business Development and is accurate as of June 2026. Service offerings, team structures, and pricing are indicative and subject to confirmation during the scoping process. Pricing shown for managed services tiers is in USD and excludes applicable taxes.